

Copy-Ready Application Texts

First-person, evidence-bounded text blocks

Short profile

I am an independent AI systems builder focused on local agent infrastructure, secure tool use, memory, model routing and evidence-bounded research.

Medium profile

I have built a personal AI operating layer around local agents, model orchestration, remote command loops, memory systems, voice control and experimental evaluation. The work is not presented as a finished commercial product; it is presented as evidence that I can independently build and reason about the real infrastructure problems behind secure agent systems.

Long profile

My work started from practical need: I wanted AI systems that could help me from my phone, speak with me, remember context, use different models intelligently, and operate local tools without becoming unsafe. That led to a broad but connected architecture: WhatsApp/remote ingress, Dual-Brain interpretation and evaluation, Miguel Core as local control plane, model cascade routing, CASS and Memory Broker, a voice tool interface, and ANIMA as an experimental cognitive kernel. The strongest value is the system intuition I developed by building these pieces myself.

Why Foresight

Foresight would help turn an unusually broad personal laboratory into mentored, reproducible secure-agent evaluation work.

Project answer

I would build a secure-agent and AI-for-science proof of concept around uncertainty-aware AMR modeling, starting with synthetic/public data and explicit governance.

Notable work answer

The strongest work to mention is the local agent operating layer: remote command ingress, Dual-Brain evaluation, Miguel Core, model cascade, memory governance and voice control. I would describe it as a personal secure-agent laboratory rather than as a single app.

Growth answer

What I need now is stronger review, better benchmarks, public-safe demos, mentorship and a community that can turn breadth into rigorous research output.

Interview proof points

- Remote phone-to-agent workflow with evaluator loop.
- Local Miguel Core daemon/control layer.
- Model routing across local/free/paid/provider tiers.
- CASS and Memory Broker as memory reliability infrastructure.
- ANIMA as tested experimental cognitive kernel.
- ECSA as future project with uncertainty and governance.

How I would answer a skeptical evaluator

I would not ask them to accept the portfolio as proof of production readiness. I would say: here is the evidence that I can build, connect and reason about complex agent systems; here are the boundaries; here is what I would harden next with mentorship and external review.

Claims I should avoid

- Do not claim production-hardened security.
- Do not claim clinical validation for AMR.
- Do not claim proven artificial consciousness for ANIMA.
- Do not claim granted patent status without verified legal evidence.
- Do not lead with Quantum/EXON in the Fellowship context.