

Remote agent command loop

Individual printable topic report.

Generated 2026-06-17

Why I built it

A phone should be able to issue work to a home computer without turning the system into an unsafe direct shell. The design had to translate casual WhatsApp messages or audio into structured work orders, run them through a tool-enabled agent side, judge completion, and only then report back.

Architecture and evidence

Implemented across OpenClaw watcher, MiguelAgent, EXOVIUM core, WA task manager, cognition routing, and architecture notes. Source notes document immediate ACK behavior, quality-loop thresholding, model routing, audio transcription, and two successful end-to-end tests.

Claim boundary: Prototype/evolved system. Some EXOVIUM pieces were later superseded by Miguel Core. Private channel identifiers and credentials are excluded.

- WhatsApp message or audio arrives through OpenClaw.
- A watcher normalizes inbound events, handles media, deduplicates repeated messages, and dispatches work.
- MiguelAgent interprets the request into goal, complexity, work order, success criteria, and model hint.
- The execution side uses EXOVIUM/Miguel tooling and model routing.
- A quality loop evaluates the result and can ask for an improved work order before replying.
- The final response is formatted and sent back through WhatsApp.

What I can show with this

- Independent design of a phone-to-computer agent workflow before relying on a managed product workflow.
- Separation of front-agent understanding/evaluation from back-agent execution.
- Ability to turn vague mobile input into structured work orders with success criteria.
- Awareness that autonomous work needs progress, interruption, deduplication, and post-work evaluation.
- Remote control is an authority problem, not just an interface problem.
- Acknowledge fast, execute slowly, and report only after evaluation.
- Audio, chat, browser events, model output, and tool execution all need different failure handling.
- Credentials and channel identifiers must be isolated from public documentation.
- Replace private ad-hoc credentials with explicit scoped tokens.
- Add command-class permissions and confirmation thresholds.
- Store execution traces in a reviewer-safe audit log.
- Retest the end-to-end path after Miguel Core consolidation.

